

Part IV

Enterprise Beans

Part IV explores Enterprise JavaBeans components.

This part contains the following chapters:

- Chapter 22, Enterprise Beans
- Chapter 23, Getting Started with Enterprise Beans
- Chapter 24, Running the Enterprise Bean Examples
- Chapter 25, A Message-Driven Bean Example
- Chapter 26, Using the Embedded Enterprise Bean Container
- Chapter 27, Using Asynchronous Method Invocation in Session Beans

Chapter 22

Enterprise Beans

Enterprise **beans** are Java EE components that implement Enterprise JavaBeans (**EJB**) technology.

Enterprise **beans** run in the **EJB** container, a runtime environment within the GlassFish Server (see Container Types).

Although transparent to the application developer, the **EJB** container provides system-level services, such as transactions and security, to its enterprise **beans**.

These services enable you to quickly build and deploy enterprise **beans**, which form the core of transactional Java EE applications.

The following topics are addressed here:

- . What Is an Enterprise Bean?
- . What Is a Session Bean?
- . What Is a Message-Driven Bean?
- . Accessing Enterprise Beans
- . The Contents of an Enterprise Bean
- . Naming Conventions for Enterprise Beans
- . The Lifecycles of Enterprise Beans
- . Further Information about Enterprise Beans

What Is an Enterprise Bean?

Written in the Java programming language, an enterprise **bean** is a server-side component that encapsulates the **business** logic of an application.

The **business** logic is the code that fulfills the purpose of the application.

In an inventory control application, for example, the enterprise **beans** might implement the **business** logic in methods called **checkInventoryLevel** and **orderProduct**.

By invoking these 32-bit methods, clients can access the inventory services provided by the application.

Benefits of Enterprise Beans

For several reasons, enterprise beans simplify the development of large, distributed applications.

First, because the EJB container provides system-level services to enterprise beans, the bean developer can concentrate on solving business problems.

The **EJB** container, rather than the **bean** developer, is responsible for system-level services, such as transaction management and security authorization.

Second, because the **beans** rather than the clients contain the application's **business** logic, the client developer can focus on the presentation of the client.

The client developer does not have to code the routines that implement **business** rules or access **databases**.

As a result, the clients are thinner, a benefit that is particularly important for clients that run on small devices.

Third, because enterprise beans are portable components, the application assembler can build new applications from existing beans.

Provided that they use the standard APIs, these applications can run on any compliant Java EE server.

When to Use Enterprise Beans

You should consider using enterprise beans if your application has any of the following requirements.

- . The application must be scalable.

To accommodate a growing number of users, you may need to distribute an application's components across multiple machines.

Not only can the enterprise beans of an application run on different machines, but also their location will remain transparent to the clients.

- . Transactions must ensure **data integrity**.

Enterprise **beans** support transactions, the mechanisms that **manage** the concurrent access of shared **objects**.

- . The application will have a variety of clients.

With only a few lines of code, remote clients can easily locate enterprise **beans**.

These clients can be thin, various, and numerous.

Types of Enterprise Beans

Table 22-1 summarizes the two types of enterprise beans.

The following sections discuss each type in more detail.

Table 22-1 Enterprise Bean Types

Enterprise Bean Type	Purpose
Session	Performs a task for a client; optionally, may implement a web service
Message-driven	Acts as a listener for a particular messaging type, such as the Java Message Service API

What Is a Session Bean?

A **session bean** encapsulates **business** logic that can be invoked programmatically by a client over local, remote, or web service client views.

To access an application that is deployed on the server, the client invokes the session **bean's** methods.

The session **bean** performs work for its client, shielding it **from** complexity by executing **business** tasks inside the server.

A session **bean** is not persistent.

(That is, its **data** is not saved to a **database**.)

For code samples, see Chapter 24, Running the Enterprise **Bean** Examples.

Types of Session Beans

Session beans are of three types: stateful, stateless, and singleton.

Stateful Session Beans

The state of an **object** consists of the values of its instance variables.

In a **stateful session bean**, the instance variables represent the state of a unique client/**bean** session.

Because the client **interacts** (“talks”) with its **bean**, this state is often called the **conversational state**.

As its name suggests, a session **bean** is similar to an **interactive session**.

A session **bean** is not shared; it can have only one client, in the same way that an **interactive session** can have only one user.

When the client terminates, its session **bean** appears to terminate and is no longer associated with the client.

The state is retained for the duration of the client/**bean** session.

If the client removes the **bean**, the session ends and the state disappears.

This transient nature of the state is not a problem, however, because when the conversation between the client and the bean ends, there is no need to retain the state.

Stateless Session Beans

A **stateless session bean** does not maintain a conversational state with the client.

When a client invokes the methods of a stateless **bean**, the **bean's** instance variables may contain a state specific to that client but only for the duration of the invocation.

When the method is finished, the client-specific state should not be retained.

Clients may, however, change the state of instance variables in pooled stateless beans, and this state is held over to the next invocation of the pooled stateless bean.

Except during method invocation, all instances of a stateless bean are equivalent, allowing the EJB container to assign an instance to any client.

That is, the state of a stateless session bean should apply across all clients.

Because they can support multiple clients, stateless session beans can offer better scalability for applications that require large numbers of clients.

Typically, an application requires fewer stateless session beans than stateful session beans to support the same number of clients.

A stateless session bean can implement a web service, but a stateful session bean cannot.

Singleton Session Beans

A **singleton session bean** is instantiated once per application and exists for the **lifecycle** of the application.

Singleton session **beans** are **designed** for circumstances in which a single enterprise **bean** instance is shared across and concurrently accessed by clients.

Singleton session **beans** offer similar functionality to stateless session **beans** but differ **from** them in that there is only one singleton session **bean** per application, as opposed to a **pool** of stateless session **beans**, any of which may respond to a client request.

Like stateless session **beans**, singleton session **beans** can implement web service endpoints.

Singleton session **beans** **maintain** their state between client invocations but are not **required** to **maintain** their state across server crashes or shutdowns.

Applications that **use** a singleton session **bean** may specify that the singleton should be instantiated upon application startup, which allows the singleton to perform initialization tasks for the application.

The singleton may perform cleanup tasks on application shutdown as well, because the singleton will operate throughout the lifecycle of the application.

When to Use Session Beans

Stateful session **beans** are appropriate if any of the following conditions are true.

- . The **bean**'s state represents the **interaction** between the **bean** and a specific client.
- . The **bean** needs to hold information about the client across method invocations.

- . The **bean** mediates between the client and the other components of the application, presenting a simplified view to the client.
- . Behind the scenes, the **bean** manages the work flow of several enterprise **beans**.

To improve performance, you might choose a stateless session **bean** if it has any of these traits.

- . The **bean**'s state has no **data** for a specific client.
- . In a single method invocation, the **bean** performs a generic task for all clients.

For example, you might use a stateless session **bean** to send an email that confirms an online order.

- . The **bean** implements a web service.

Singleton session **beans** are appropriate in the following circumstances.

- . State needs to be shared across the application.
- . A single enterprise **bean** needs to be accessed by multiple threads concurrently.
- . The application needs an enterprise **bean** to perform tasks upon application startup and shutdown.
- . The **bean** implements a web service.

What Is a Message-Driven Bean?

A **message-driven bean** is an enterprise **bean** that allows Java EE applications to **process** messages asynchronously.

This type of **bean** normally acts as a JMS message listener, which is similar to an event listener but receives JMS messages instead of events.

The messages can be sent by any Java EE component (an application client, another enterprise bean, or a web component) or by a JMS application or system that does not use Java EE technology.

Message-driven beans can process JMS messages or other kinds of messages.

What Makes Message-Driven Beans Different from Session Beans?

The most visible difference between message-driven beans and session beans is that clients do not access message-driven beans through interfaces.

Interfaces are described in the section Accessing Enterprise Beans.

Unlike a session **bean**, a message-driven **bean** has only a **bean class**.

In several respects, a message-driven **bean** resembles a stateless session **bean**.

- . A message-driven **bean**'s instances retain no **data** or conversational state for a specific client.

- . All instances of a message-driven bean are equivalent, allowing the EJB container to assign a message to any message-driven bean instance.

The container can pool these instances to allow streams of messages to be processed concurrently.

- . A single message-driven bean can process messages from multiple clients.

The instance variables of the message-driven bean instance can contain some state across the handling of client messages, such as a JMS API connection, an open database connection, or an object reference to an enterprise bean object.

Client components do not locate message-driven beans and invoke methods directly on them.

Instead, a client accesses a message-driven bean through, for example, JMS by sending messages to the message destination for which the message-driven bean class is the **MessageListener**.

You assign a message-driven **bean**'s destination during deployment by using GlassFish Server resources.

Message-driven **beans** have the following characteristics.

- . They execute upon receipt of a single client message.
- . They are invoked asynchronously.
- . They are relatively short-lived.

- . They do not represent directly shared **data** in the **database**, but they can access and update this **data**.
- . They can be transaction-aware.
- . They are stateless.

When a message arrives, the container calls the message-driven **bean**'s **onMessage** method to **process** the message.

The **onMessage** method normally casts the message to one of the five JMS message types and handles it in accordance with the application's **business** logic.

The **onMessage** method can call helper methods or can invoke a session **bean** to **process** the information in the message or to store it in a **database**.

A message can be delivered to a message-driven bean within a transaction context, so all operations within the `onMessage` method are part of a single transaction.

If message processing is rolled back, the message will be redelivered.

For more information, see Chapter 43,
Transactions.

When to Use Message-Driven Beans

Session beans allow you to send JMS messages and to receive them synchronously but not asynchronously.

To avoid tying up server resources, do not to use blocking synchronous receives in a server-side component; in general, JMS messages should not be sent or received synchronously.

To receive messages asynchronously, use a message-driven bean.

Accessing Enterprise Beans

Note - The material in this section applies only to session beans and not to message-driven beans.

Because they have a different programming model, message-driven beans do not have interfaces or no-interface views that define client access.

Clients access enterprise **beans** either through a **no-interface view** or through a **business interface**.

A **no-interface view** of an enterprise **bean** exposes the public methods of the enterprise **bean** implementation **class** to clients.

Clients using the no-**interface** view of an enterprise **bean** may invoke any public methods in the enterprise **bean** implementation **class** or any super**classes** of the implementation **class**.

A **business interface** is a standard Java programming language **interface** that contains the **business** methods of the enterprise **bean**.

A client can access a session **bean** only through the methods defined in the **bean's business interface** or through the public methods of an enterprise **bean** that has a no-**interface** view.

The **business interface** or no-**interface** view defines the client's view of an enterprise **bean**.

All other aspects of the enterprise **bean** (method implementations and deployment settings) are hidden **from** the client.

Well-**designed** **interfaces** and **no-interfaces** views simplify the development and **maintenance** of Java EE applications.

Not only do clean **interfaces** and **no-interfaces** views shield the clients **from** any complexities in the **EJB** tier, but they also allow the enterprise **beans** to change **internally** without affecting the clients.

For example, if you change the implementation of a session **bean** **business** method, you won't have to alter the client code.

But if you were to change the method definitions in the **interfaces**, you might have to modify the client code as well.

Therefore, it is important that you **design** the **interfaces** and **no-interface** views carefully to isolate your clients **from** possible changes in the enterprise **beans**.

Session **beans** can have more than one **business interface**.

Session **beans** should, but are not **required** to, implement their **business interface** or **interfaces**.

Using Enterprise Beans in Clients

The client of an enterprise **bean** obtains a reference to an instance of an enterprise **bean** through either **dependency injection**, using Java programming language annotations, or **JNDI lookup**, using the Java Naming and Directory Interface syntax to find the enterprise **bean** instance.

Dependency injection is the simplest way of obtaining an enterprise **bean** reference.

Clients that run within a Java EE server-**managed** environment, JavaServer Faces web applications, JAX-**RS** web services, other enterprise **beans**, or Java EE application clients, support dependency injection using the **javax.ejb.EJB** annotation.

Applications that run outside a Java EE server-managed environment, such as Java SE applications, must perform an explicit lookup.

JNDI supports a global syntax for identifying Java EE components to simplify this explicit lookup.

Portable JNDI Syntax

Three JNDI namespaces are used for portable JNDI lookups: `java:global`, `java:module`, and `java:app`.

- . The `java:global` JNDI namespace is the portable way of finding remote enterprise beans using JNDI lookups.

JNDI addresses are of the following form:

```
java:global  
[/application name]  
/module name/enterprise bean  
name[/interface name]
```

Application name and module name default to the name of the application and module minus the file extension.

Application names are **required** only if the application is packaged within an EAR.

The **interface** name is **required** only if the enterprise **bean** implements more than one **business interface**.

- . The **java:module** namespace is **used** to **look** up local enterprise **beans** within the same module.

JNDI addresses using the `java:module` namespace are of the following form:

`java:module/enterprise bean
name/[interface name]`

The `interface` name is `required` only if the enterprise `bean` implements more than one business `interface`.

- The `java:app` namespace is used to look up local enterprise beans packaged within the same application.

That is, the enterprise bean is packaged within an EAR file containing multiple Java EE modules.

JNDI addresses using the `java:app` namespace are of the following form:


```
java:app[/module name]  
/enterprise bean name  
[/interface name]
```

The module name is optional.

The **interface** name is **required** only if the enterprise **bean** implements more than one **business interface**.

For example, if an enterprise bean, **MyBean**, is packaged within the web application archive **myApp.war**, the module name is **myApp**.

The portable JNDI name is

java:module/MyBean An equivalent JNDI name using the **java:global** namespace is **java:global/myApp/MyBean**.

Deciding on Remote or Local Access

When you **design** a Java EE application, one of the first decisions you make is the type of client access allowed by the enterprise **beans**: remote, local, or web service.

Whether to allow local or remote access depends on the following factors.

- **Tight or loose coupling of related beans:**
Tightly coupled beans depend on one another.

For example, if a session bean that processes sales orders calls a session bean that emails a confirmation message to the customer, these beans are tightly coupled.

Tightly coupled beans are good candidates for local access.

Because they fit together as a logical unit, they typically call each other often and would benefit from the increased performance that is possible with local access.

- . **Type of client:** If an enterprise bean is accessed by application clients, it should allow remote access.

In a production environment, these clients almost always run on machines other than those on which the GlassFish Server is running.

If an enterprise **bean**'s clients are web components or other enterprise **beans**, the type of access depends on how you want to distribute your components.

- **Component distribution:** Java EE applications are scalable because their server-side components can be distributed across multiple machines.

In a distributed application, for example, the server that the web components run on may not be the one on which the enterprise beans they access are deployed.

In this distributed scenario, the enterprise beans should allow remote access.

- . **Performance:** Owing to such factors as network latency, remote calls may be slower than local calls.

On the other hand, if you distribute components among different servers, you may improve the application's overall performance.

**Both of these statements are generalizations;
performance can vary in different operational
environments.**

**Nevertheless, you should keep in mind how
your application **design** might affect
performance.**

If you aren't sure which type of access an enterprise bean should have, choose remote access.

This decision gives you more flexibility.

In the future, you can distribute your components to accommodate the growing demands on your application.

Although it is uncommon, it is possible for an enterprise bean to allow both remote and local access.

If this is the case, either the business interface of the bean must be explicitly designated as a business interface by being decorated with the `@Remote` or `@Local` annotations, or the bean class must explicitly designate the business interfaces by using the `@Remote` and `@Local` annotations.

The same **business interface** cannot be both a local and a remote **business interface**.

Local Clients

A local client has these characteristics.

- . It must run in the same application as the enterprise bean it accesses.**
- . It can be a web component or another enterprise bean.**

- . To the local client, the location of the enterprise **bean** it accesses is not transparent.

The no-**interface** view of an enterprise **bean** is a local view.

The public methods of the enterprise **bean** implementation **class** are exposed to local clients that access the no-**interface** view of the enterprise **bean**.

Enterprise **beans** that use the no-**interface** view do not implement a **business interface**.

The **local business interface** defines the **bean's** **business** and lifecycle methods.

If the **bean's business interface** is not decorated with **@Local** or **@Remote**, and if the **bean class** does not specify the **interface** using **@Local** or **@Remote**, the **business interface** is by default a **local interface**.

To build an enterprise **bean** that allows only local access, you may, but are not **required to**, do one of the following:

- Create an enterprise **bean** implementation **class** that does not implement a **business interface**, indicating that the **bean** exposes a **no-interface** view to clients.

For example:

```
@Session  
public class MyBean { ... }
```

- . Annotate the **business interface** of the enterprise **bean** as a **@Local interface**.

For example:

```
@Local  
public interface InterfaceName  
{ ... }
```

- Specify the **interface** by decorating the **bean class** with **@Local** and specify the **interface name**.

For example:

```
@Local(InterfaceName.class)  
public class BeanName implements  
InterfaceName { ... }
```

Accessing Local Enterprise Beans

Using the No-Interface View

Client access to an enterprise bean that exposes a local, no-interface view is accomplished through either dependency injection or JNDI lookup.

- . To obtain a reference to the no-interface view of an enterprise bean through dependency injection, use the `javax.ejb.EJB` annotation and specify the enterprise bean's implementation class:

- . `@EJB`
`ExampleBean exampleBean;`

- . To obtain a reference to the no-interface view of an enterprise bean through JNDI lookup, use the `javax.naming.InitialContext` interface's `lookup` method:

```
ExampleBean exampleBean =  
(ExampleBean)  
InitialContext.lookup  
("java:module/ExampleBean");
```

Clients **do not** use the **new** operator to obtain a **new** instance of an enterprise **bean** that uses a **no-interface** view.

Accessing Local Enterprise Beans That Implement Business Interfaces

Client access to enterprise **beans** that implement local **business interfaces** is accomplished through either dependency injection or JNDI **lookup**.

- . To obtain a reference to the local **business interface** of an enterprise **bean** through dependency injection, use the **javax.ejb.EJB** annotation and specify the enterprise **bean's** local **business interface** name:

- . **@EJB**
Example example;

- . To obtain a reference to a local business interface of an enterprise bean through JNDI lookup, use the `javax.naming.InitialContext` interface's `lookup` method:

```
ExampleLocal example =  
(ExampleLocal)  
InitialContext.lookup  
("java:module/ExampleLocal");
```

Remote Clients

A remote client of an enterprise **bean** has the following traits.

- . It can run on a different machine and a different JVM **from** the enterprise **bean** it accesses.

(It is not **required** to run on a different JVM.)

- . It can be a web component, an application client, or another enterprise bean.
- . To a remote client, the location of the enterprise bean is transparent.
- . The enterprise bean must implement a business interface.

That is, remote clients **may not** access an enterprise bean through a no-interface view.

To create an enterprise bean that allows remote access, you must either

- . Decorate the business interface of the enterprise bean with the `@Remote` annotation:

```
@Remote  
public interface InterfaceName  
{ ... }
```

- Decorate the bean class with `@Remote`, specifying the business interface or interfaces:

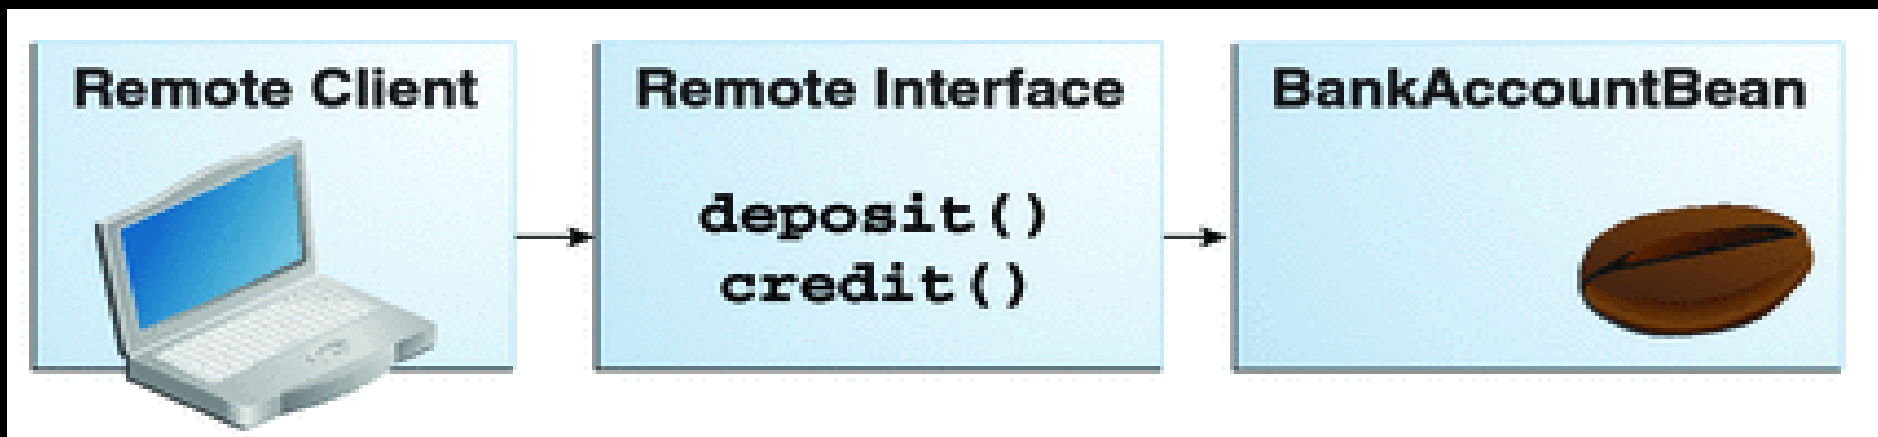
```
@Remote (InterfaceName.class)  
public class BeanName implements  
InterfaceName { ... }
```

The **remote interface** defines the **business** and **lifecycle** methods that are specific to the **bean**.

For example, the remote **interface** of a **bean** named **BankAccountBean** might have **business** methods named **deposit** and **credit**.

Figure 22-1 shows how the **interface** controls the client's view of an enterprise **bean**.

Figure 22-1 **Interfaces** for an Enterprise **Bean** with Remote Access



Client access to an enterprise **bean** that implements a remote **business interface** is accomplished through either dependency injection or JNDI **lookup**.

- . To obtain a reference to the remote **business interface** of an enterprise **bean** through dependency injection, use the **javax.ejb.EJB** annotation and specify the enterprise **bean**'s remote **business interface** name:


```
@EJB
```

```
Example example;
```

- . To obtain a reference to a remote **business interface** of an enterprise **bean** through JNDI **lookup**, use the **javax.naming.InitialContext** interface's **lookup** method:

```
ExampleRemote example =  
    (ExampleRemote)  
    InitialContext.lookup  
    ("java:global/myApp/ExampleRemote");
```

Web Service Clients

A web service client can access a Java EE application in two ways.

First, the client can access a web service created with JAX-WS.

(For more information on JAX-WS, see Chapter 18, Building Web Services with JAX-WS.)

Second, a web service client can invoke the business methods of a stateless session bean.

Message beans cannot be accessed by web service clients.

Provided that it **uses** the correct protocols (**SOAP**, **HTTP**, **WSDL**), any web service client can access a stateless session **bean**, whether or not the client is written in the Java programming language.

The client doesn't even “know” what technology implements the service: stateless session **bean**, **JAX-WS**, or some other technology.

In addition, enterprise **beans** and web components can be clients of web services.

This flexibility enables you to **integrate** Java EE applications with web services.

A web service client accesses a stateless session **bean** through the **bean's** web service endpoint implementation **class**.

By default, all public methods in the **bean class** are accessible to web service clients.

The **@WebMethod** annotation may be used to customize the **behavior** of web service methods.

If the **@WebMethod** annotation is used to decorate the **bean class's** methods, only those methods decorated with **@WebMethod** are exposed to web service clients.

For a code sample, see [A Web Service Example:](#)
`helloservice`.

Method Parameters and Access

The type of access affects the parameters of the **bean** methods that are called by clients.

The following sections apply not only to method parameters but also to method return values.

Isolation

The parameters of remote calls are more isolated than those of local calls.

With remote calls, the client and the **bean** operate on different copies of a parameter **object**.

If the client changes the value of the **object**, the value of the copy in the **bean** does not change.

This layer of isolation can help protect the bean if the client accidentally modifies the data.

In a local call, both the client and the bean can modify the same parameter object.

In general, you should not rely on this side effect of local calls.

Perhaps someday you will want to distribute your components, replacing the local calls with remote ones.

As with remote clients, web service clients operate on different copies of parameters than does the **bean** that implements the web service.

Granularity of Accessed Data

Because remote calls are likely to be slower than local calls, the parameters in remote methods should be relatively coarse-grained.

A coarse-grained **object** contains more **data** than a fine-grained one, so fewer access calls are required.

For the same reason, the parameters of the methods called by web service clients should also be coarse-grained.

The Contents of an Enterprise Bean

To develop an enterprise **bean**, you must provide the following files:

- . **Enterprise bean class**: Implements the **business** methods of the enterprise **bean** and any **lifecycle** callback methods.

- **Business interfaces:** Define the business methods implemented by the enterprise bean class.

A business interface is not required if the enterprise bean exposes a local, no-interface view.

- **Helper classes:** Other classes needed by the enterprise bean class, such as exception and utility classes.

Package the programming artifacts in the preceding list either **into** an **EJB** JAR file (a stand-alone module that stores the enterprise bean) or within a web application archive (WAR) module.

Packaging Enterprise Beans in EJB JAR Modules

An **EJB** JAR file is **portable** and can be used for various applications.

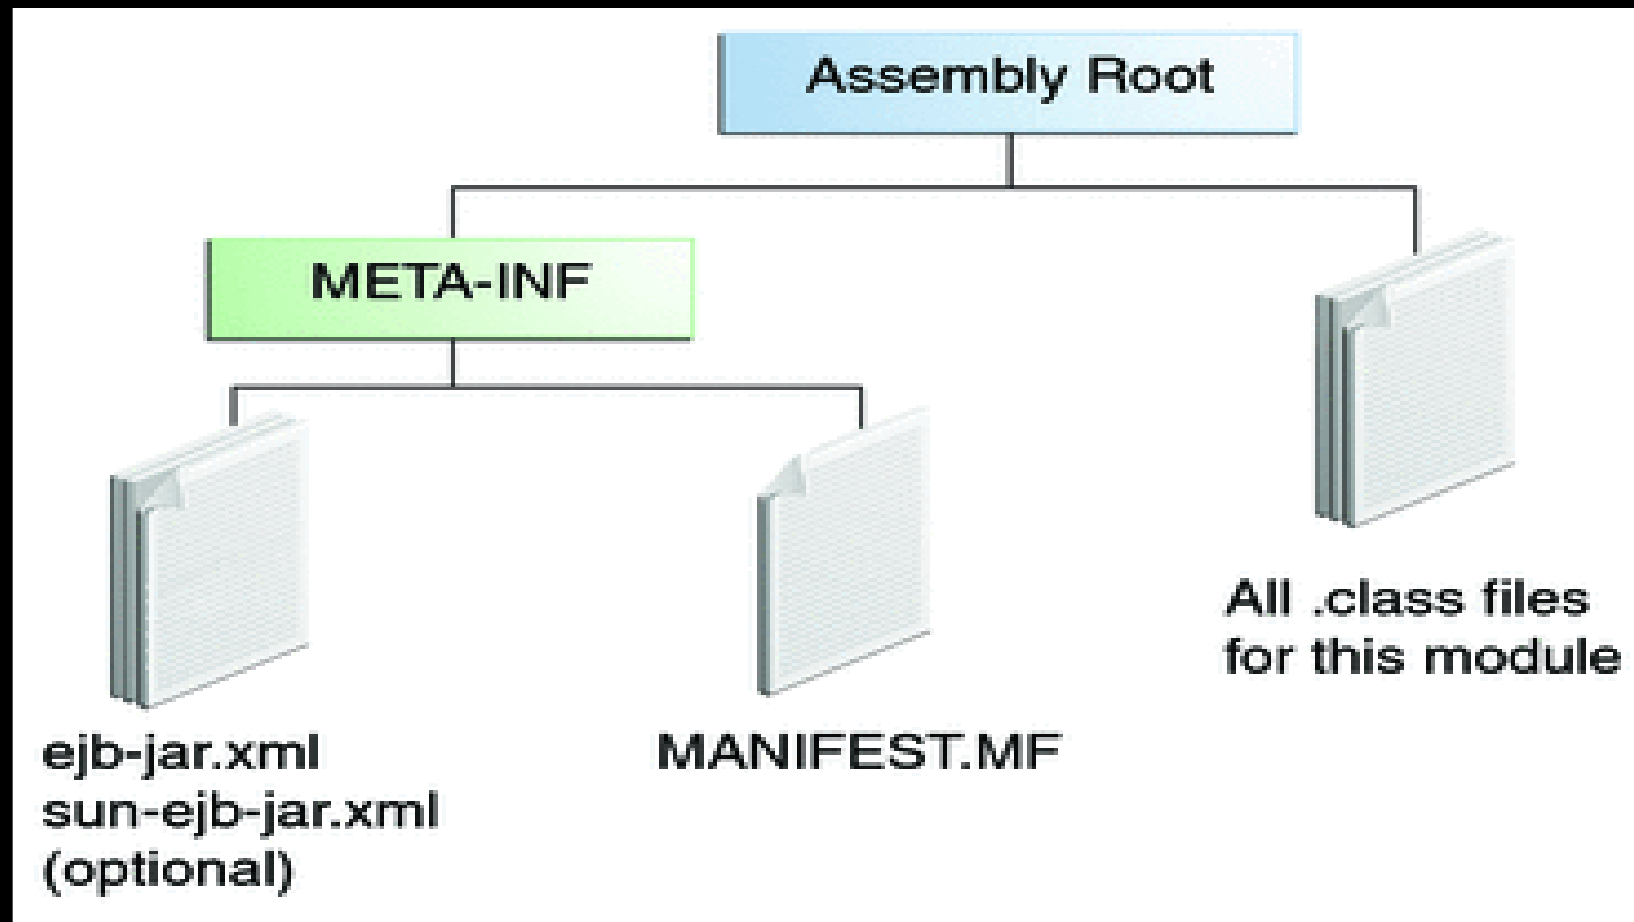
To assemble a Java EE application, package one or more modules, such as **EJB** JAR files, into an EAR file, the archive file that holds the application.

When deploying the EAR file that contains the enterprise bean's **EJB** JAR file, you also deploy the enterprise bean to the GlassFish Server.

You can also deploy an **EJB** JAR that is not contained in an EAR file.

Figure 22-2 shows the contents of an **EJB JAR** file.

Figure 22-2 Structure of an Enterprise Bean JAR



Packaging Enterprise Beans in WAR Modules

Enterprise **beans** often provide the **business** logic of a web application.

In these **cases**, packaging the enterprise **bean** within the web application's WAR module simplifies deployment and application organization.

Enterprise **beans** may be packaged within a WAR module as Java programming language **class** files or within a JAR file that is bundled within the WAR module.

To include enterprise **bean class** files in a WAR module, the **class** files should be in the **WEB-INF/classes** directory.

To include a JAR file that contains enterprise beans in a WAR module, add the JAR to the **WEB-INF/lib** directory of the WAR module.

WAR modules that contain enterprise beans do not require an **ejb-jar.xml** deployment descriptor.

If the application uses `ejb-jar.xml`, it must be located in the WAR module's `WEB-INF` directory.

JAR files that contain enterprise **bean classes** packaged within a WAR module are not considered **EJB** JAR files, even if the bundled JAR file conforms to the format of an **EJB** JAR file.

The enterprise **beans** contained within the JAR file are semantically equivalent to enterprise **beans** located in the WAR module's

WEB-INF/classes directory, and the environment namespace of all the enterprise **beans** are scoped to the WAR module.

For example, suppose that a web application consists of a shopping cart enterprise bean, a credit card processing enterprise bean, and a Java servlet front end.

The shopping cart bean exposes a local, no-interface view and is defined as follows:

```
package com.example.cart;  
@Stateless  
public class CartBean { ... }
```

The credit card **processing bean** is packaged within its own JAR file, **cc.jar**, exposes a local, **no-interface view**, and is defined as follows:

```
package com.example.cc;  
@Stateless  
public class CreditCardBean{ ... }
```

The servlet, `com.example.web.StoreServlet`, handles the web front end and uses both `CartBean` and `CreditCardBean`.

The WAR module layout for this application looks as follows:

```
WEB-INF/classes/com/example/cart/  
CartBean.class
```

WEB-INF/classes/com/example/web/
StoreServlet

WEB-INF/lib/cc.jar

WEB-INF/ejb-jar.xml

WEB-INF/web.xml

Naming Conventions for Enterprise Beans

Because enterprise **beans** are composed of multiple parts, it's useful to follow a naming convention for your applications.

Table 22-2 summarizes the conventions for the example **beans** in this tutorial.

Table 22-2 Naming Conventions for Enterprise Beans

Item	Syntax	Example
Enterprise bean name	<i>name</i> Bean	AccountBean
Enterprise bean class	<i>name</i> Bean	AccountBean
Business interface	<i>name</i>	Account

The Lifecycles of Enterprise Beans

An enterprise **bean** goes through various **stages** during its **lifetime**, or **lifecycle**.

Each type of enterprise **bean** (stateful session, stateless session, singleton session, or message-driven) has a different lifecycle.

The **descriptions** that follow refer to methods that are explained along with the code examples in the next two chapters.

If you are **new** to enterprise **beans**, you should skip this section and run the code examples first.

The Lifecycle of a Stateful Session Bean

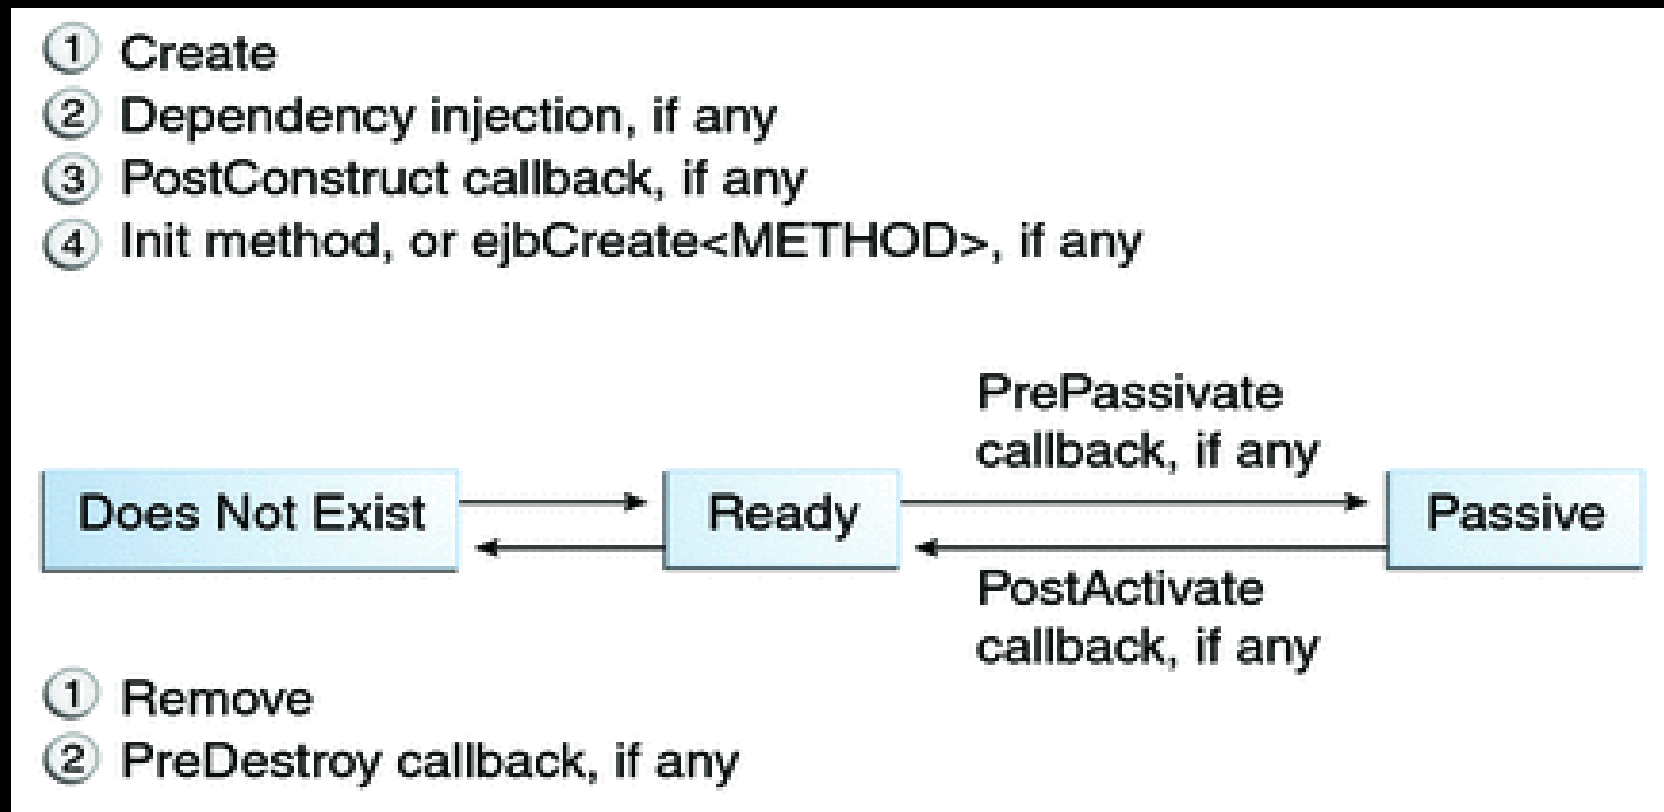
Figure 22-3 illustrates the **stages** that a session **bean** passes through during its lifetime.

The client initiates the **lifecycle** by obtaining a reference to a stateful session **bean**.

The container performs any dependency injection and then invokes the method annotated with `@PostConstruct`, if any.

The **bean** is now ready to have its **business** methods invoked by the client.

Figure 22-3 Lifecycle of a Stateful Session Bean



While in the ready stage, the EJB container may decide to deactivate, or passivate, the bean by moving it from memory to secondary storage.

(Typically, the **EJB** container uses a least-recently-used algorithm to **select** a **bean** for passivation.) The **EJB** container invokes the method annotated **@PrePassivate**, if any, immediately before passivating it.

If a client invokes a **business** method on the **bean** while it is in the passive **stage**, the **EJB** container activates the **bean**, calls the method annotated **@PostActivate**, if any, and then moves it to the ready **stage**.

At the end of the **lifecycle**, the client invokes a method annotated **@Remove**, and the **EJB** container calls the method annotated **@PreDestroy**, if any.

The **bean**'s instance is then ready for garbage collection.

Your code controls the invocation of only one lifecycle method: the method annotated **@Remove**.

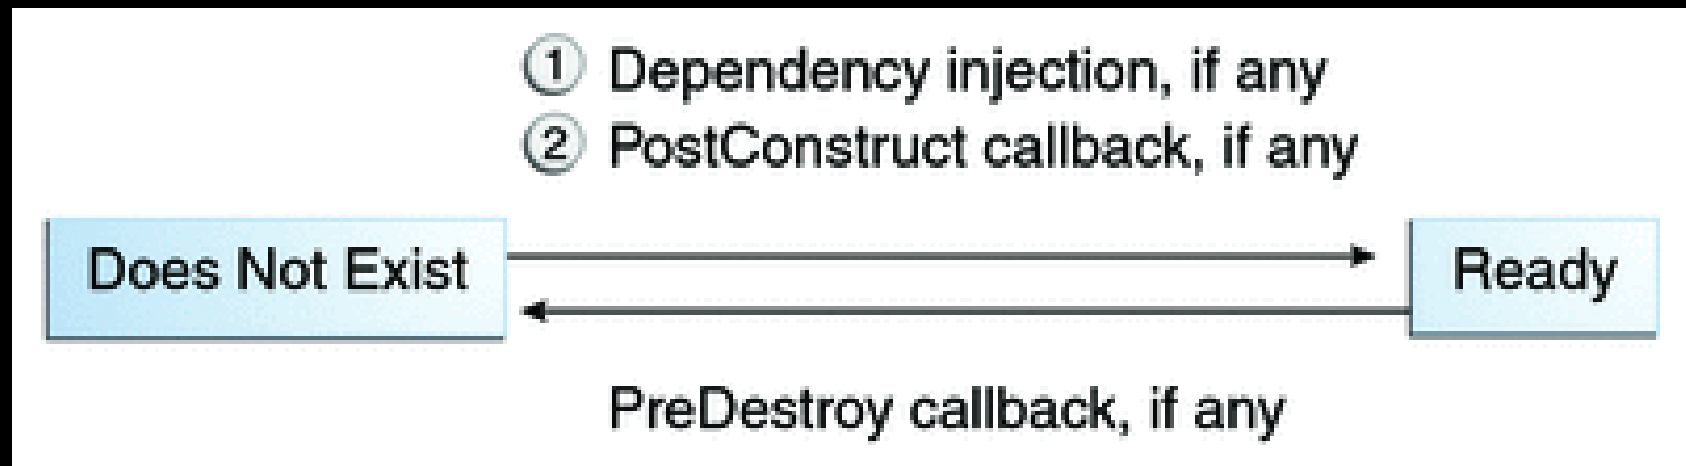
All other methods in Figure 22-3 are invoked by the **EJB** container. See Chapter 44, Resource Connections for more information.

The Lifecycle of a Stateless Session Bean

Because a stateless session **bean** is never passivated, its lifecycle has only two **stages**: nonexistent and ready for the invocation of **business** methods.

Figure 22-4 illustrates the **stages** of a stateless session **bean**.

Figure 22-4 Lifecycle of a Stateless Session Bean



The **EJB** container typically creates and maintains a pool of stateless session beans, beginning the stateless session bean's lifecycle.

The container performs any dependency injection and then invokes the method annotated **@PostConstruct**, if it exists.

The **bean** is now ready to have its **business** methods invoked by a client.

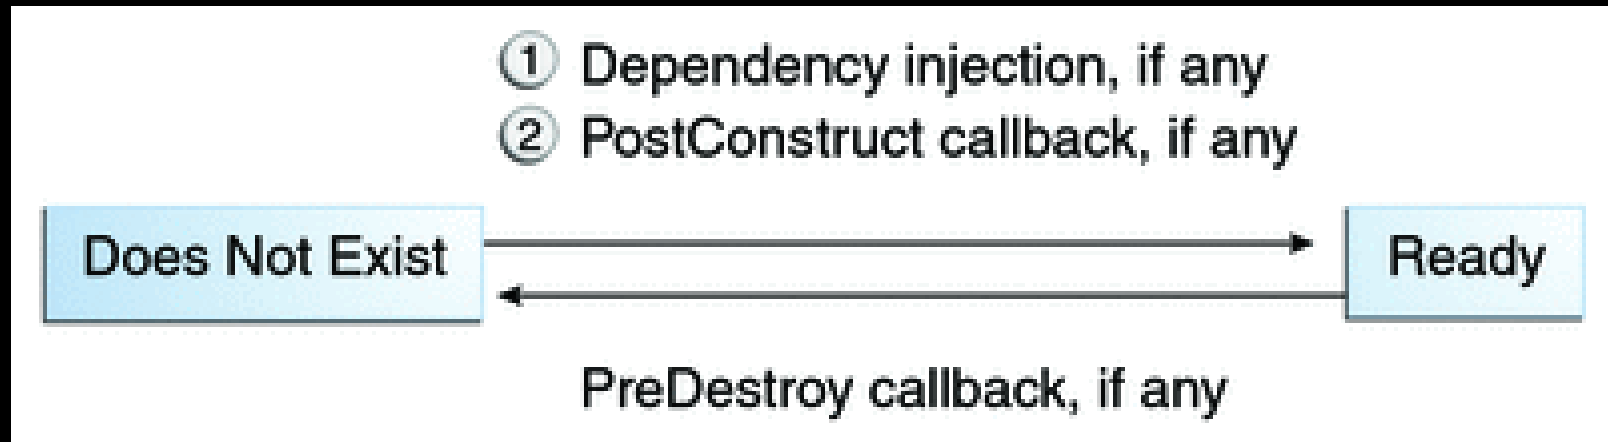
At the end of the **lifecycle**, the **EJB** container calls the method annotated **@PreDestroy**, if it exists.

The **bean**'s instance is then ready for garbage collection.

The Lifecycle of a Singleton Session Bean

Like a stateless session bean, a singleton session bean is never passivated and has only two stages, nonexistent and ready for the invocation of business methods, as shown in Figure 22-5.

Figure 22-5 Lifecycle of a Singleton Session Bean



The **EJB** container initiates the singleton session bean lifecycle by creating the singleton instance.

This occurs upon application deployment if the singleton is annotated with the `@Startup` annotation. The container performs any dependency injection and then invokes the method annotated `@PostConstruct`, if it exists.

The singleton session **bean** is now ready to have its **business** methods invoked by the client.

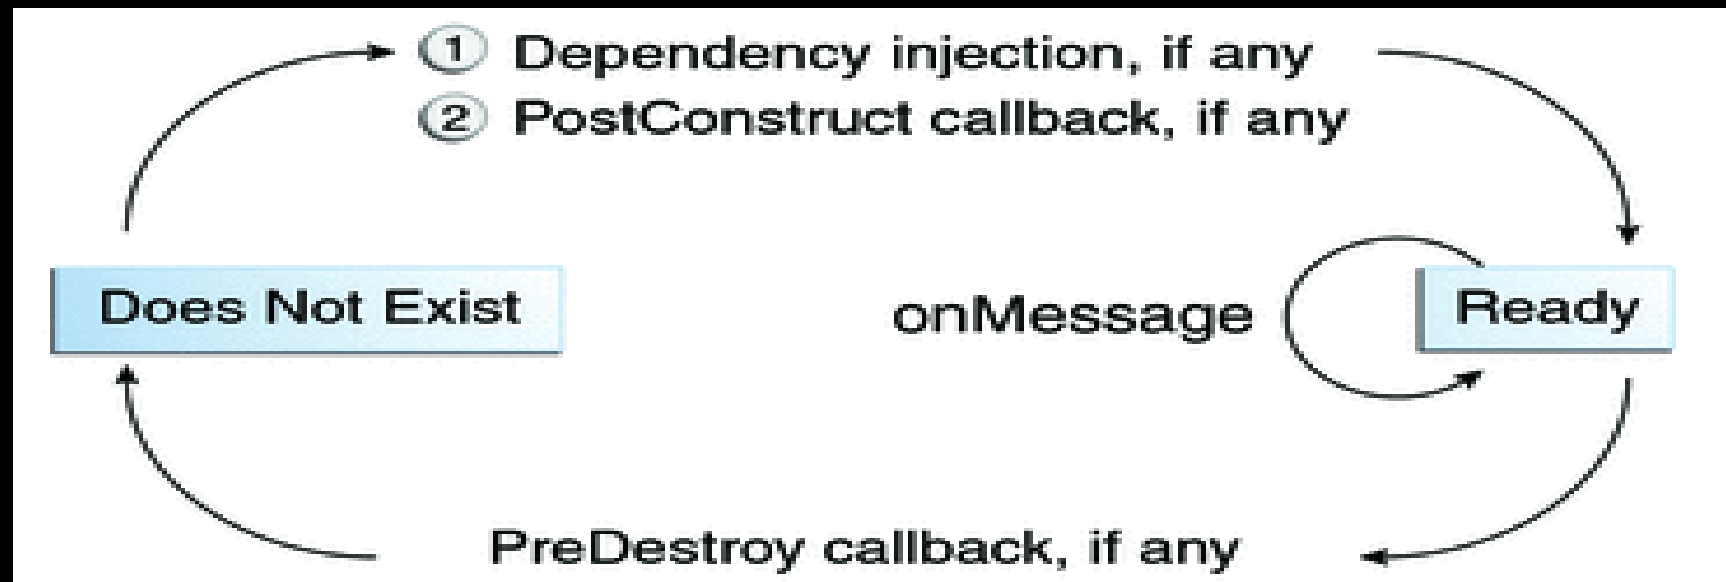
At the end of the lifecycle, the **EJB** container calls the method annotated **@PreDestroy**, if it exists.

The singleton session **bean** is now ready for garbage collection.

The Lifecycle of a Message-Driven Bean

Figure 22-6 illustrates the stages in the lifecycle of a message-driven bean.

Figure 22-6 Lifecycle of a Message-Driven Bean



The **EJB** container usually creates a pool of message-driven bean instances.

For each instance, the **EJB** container performs these tasks.

1. If the message-driven bean uses dependency injection, the container injects these references before instantiating the instance.
2. The container calls the method annotated **@PostConstruct**, if any.

Like a stateless session **bean**, a message-driven **bean** is never passivated and has only two states: nonexistent and ready to receive messages.

At the end of the **lifecycle**, the container calls the method annotated **@PreDestroy**, if any.

The **bean**'s instance is then ready for garbage collection.

Further Information about Enterprise Beans

For more information on Enterprise JavaBeans technology, see

- . Enterprise JavaBeans 3.1 specification:

<http://jcp.org/en/jsr/summary?id=318>

- . Enterprise JavaBeans web site:

<http://www.oracle.com/technetwork/java/ejb-141389.html>