

Improving the Performance of Java Persistence API Applications By Setting a Second-Level Cache

This chapter explains how to modify the second-level cache mode settings to improve the performance of applications that use the Java Persistence API.

The following topics are addressed here:

- . Overview of the Second-Level Cache
- . Specifying the Cache Mode Settings to Improve Performance

Overview of the Second-Level Cache

A **second-level cache** is a local store of entity **data** managed by the persistence provider to improve application performance.

A second-level cache helps improve performance by avoiding expensive **database** calls, keeping the entity **data** local to the application.

A second-level cache is typically transparent to the application, as it is managed by the persistence provider and underlies the persistence context of an application.

That is, the application reads and commits data through the normal entity manager operations without knowing about the cache.

Note - Persistence providers are not required to support a second-level cache.

Portable applications should not rely on support by persistence providers for a second-level cache.

The second-level cache for a persistence unit may be configured to one of several second-level cache modes.

The following cache mode settings are defined by the Java Persistence **API**:

Table 38-1 Cache Mode Settings for the Second-Level Cache

Cache Mode Setting	Description
ALL	All entity data is stored in the second-level cache for this persistence unit.
NONE	No data is cached in the persistence unit. The persistence provider must not cache any data .

ENABLE_SELECTIVE	Enable caching for entities that have been explicitly set with the @Cacheable annotation.
DISABLE_SELECTIVE	Enable caching for all entities except those that have been explicitly set with the @Cacheable(false) annotation.
UNSPECIFIED	<p>The caching behavior for the persistence unit is undefined.</p> <p>The persistence provider's default caching behavior will be used.</p>

One consequence of using a second-level cache in an application is that the underlying **data** may have changed in the **database tables**, but the value in the cache has not, a circumstance called a **stale read**.

Stale reads may be avoided by changing the second-level cache to one of the cache mode settings, controlling which entities may be cached (described in Controlling Whether Entities May Be Cached),

or changing the cache's retrieval or store modes
(described in Setting the Cache Retrieval and
Store Modes).

Which strategies best avoid stale reads are
application dependent.

Controlling Whether Entities May Be Cached

The `javax.persistence.Cacheable` annotation is used to specify that an entity class, and any subclasses, may be cached when using the `ENABLE_SELECTIVE` or `DISABLE_SELECTIVE` cache modes.

Subclasses may override the `@Cacheable` setting by adding a `@Cacheable` annotation and changing the value.

To specify that an entity may be cached, add a `Cacheable` annotation at the class level:

```
@Cacheable  
@Entity  
public class Person { ... }
```

By default, the `@Cacheable` annotation is `true`.

The following example is equivalent:

```
@Cacheable(true)
@Entity
public class Person{...}
```

To specify that an entity must not be cached, add a `@Cacheable` annotation and set it to `false`:

```
@Cacheable(false)
@Entity
public class OrderStatus {...}
```

When the **ENABLE_SELECTIVE** cache mode is set, the persistence provider will cache any entities that have a **@Cacheable(true)** annotation and any subclasses of that entity that have not been overridden.

The persistence provider will not cache entities that have **@Cacheable(false)** or have no **@Cacheable** annotation.

That is, the **ENABLE_SELECTIVE** mode will only cache entities that have been explicitly marked for the cache using the **@Cacheable** annotation.

When the **DISABLE_SELECTIVE** cache mode is set, the persistence provider will cache any entities that **do not** have a **@Cacheable(false)** annotation.

Entities that do not have a `@Cacheable` annotation, and entities with a `@Cacheable(true)` annotation will be cached.

That is, the `DISABLE_SELECTIVE` mode will cache all entities that have not been explicitly prevented **from** being cached.

If the cache mode is set to **UNDEFINED**, or is left unset, the behavior of entities annotated with **@Cacheable** is undefined.

If the cache mode is set to **ALL** or **NONE**, the value of the **@Cacheable** annotation is ignored by the persistence provider.

Specifying the Cache Mode Settings to Improve Performance

To adjust the cache mode settings for a persistence unit, specify one of the cache modes as the value of the `shared-cache-mode` element in the `persistence.xml` deployment descriptor:

```
<persistence-unit  
name="examplePU"  
transaction-type="JTA">  
  <provider>  
    org.eclipse.persistence.jpa.  
    PersistenceProvider  
  </provider>  
  <jta-data-source>  
    jdbc/__default  
  </jta-data-source>
```

```
<shared-cache-mode>  
DISABLE_SELECTIVE  
</shared-cache-mode>  
</persistence-unit>
```

Note - Because support for a second-level cache is not required by the Java Persistence API specification, setting the second-level cache mode in `persistence.xml` will have no affect when using a persistence provider that does not implement a second-level cache.

Alternately, the shared cache mode may be specified by setting the

`javax.persistence.sharedCache.mode`

property to one of the shared cache mode settings:

```
EntityManagerFactory emf =  
Persistence.  
createEntityManagerFactory  
("myExamplePU",  
new Properties().add  
("javax.persistence.sharedCache.mode"  
, "ENABLE_SELECTIVE"));
```

Setting the Cache Retrieval and Store Modes

If the second-level cache has been enabled for a persistence unit by setting the shared cache mode, the behavior of the second-level cache can be further modified by setting the `javax.persistence.cache.retrieveMode` and `javax.persistence.cache.storeMode` properties.

These properties may be set at the persistence context level by passing the property name and value to the `EntityManager.setProperty` method, or may be set on a per-`EntityManager` operation (`EntityManager.find` or `EntityManager.refresh`) or per-query level.

Cache Retrieval Mode

The cache retrieval mode, set by the `javax.persistence.retrieveMode` property, controls how data is read from the cache for calls to the `EntityManager.find` method and from queries.

The **retrieveMode** property can be set to one of the constants defined by the **javax.persistence.CacheRetrieveMode** enumerated type, either **USE** (the default) or **BYPASS**.

When set to **USE**, **data** is retrieved **from** the second-level cache, if available.

If the **data** is not in the cache, the persistence provider will read it **from** the **database**.

When set to **BYPASS**, the second-level cache is bypassed and a call to the **database** is made to retrieve the **data**.

Cache Store Mode

The cache store mode, set by the `javax.persistence.storeMode` property, controls how `data` is stored in the cache.

The **storeMode** property can be set to one of the constants defined by the **javax.persistence.CacheStoreMode** enumerated type, either **USE** (the default), **BYPASS**, or **REFRESH**.

When set to **USE** the cache **data** is created or updated when **data** is read **from** or committed to the **database**.

If **data** is already in the cache, setting the store mode to **USE** will not force a refresh when **data** is read **from** the **database**.

When the store mode is set to **BYPASS**, **data** read **from** or committed to the **database** is **not** inserted or updated in the cache.

That is, the cache is unchanged.

When the store mode is set to **REFRESH** the cache **data** is created or updated when **data** is read **from** or committed to the **database**, and a refresh is forced on **data** in the cache upon **database** reads.

Setting the Cache Retrieval or Store Mode

To set the cache retrieval or store mode for the persistence context, call the **EntityManager.setProperty** method with the property name and value pair:

```
EntityManager em = ...;  
em.setProperty("javax.persistence.c  
ache.storeMode", "BYPASS");
```

To set the cache retrieval or store mode when calling the `EntityManager.find` or `EntityManager.refresh` methods, first create a `Map<String, Object>` instance and add a name/value pair as follows:

```
EntityManager em = ...;  
Map<String, Object> props =  
new HashMap<String, Object>();
```

```
props.put(  
    "javax.persistence.cache.retrieveMode",  
    "BYPASS");  
String personPK = ...;  
Person person = em.find  
(Person.class, personPK, props);
```

Note - The cache retrieve mode is ignored when calling the `EntityManager.refresh` method, as calls to `refresh` always result in `data` being read `from` the `database`, not the cache.

To set the retrieval or store mode when using queries, call the `Query.setHint` or `TypedQuery.setHint` methods, depending on the type of `query`:

```
EntityManager em = ...;  
CriteriaQuery<Person> cq = ...;  
TypedQuery<Person> q =  
em.createQuery(cq);  
q.setHint(  
    "javax.persistence.cache.storeMode"  
    , "REFRESH");  
...
```

Setting the store or retrieve mode in a **query** or when calling the **EntityManager.find** or **EntityManager.refresh** methods overrides the setting of the entity **manager**.

Controlling the Second-Level Cache Programmatically

The `javax.persistence.Cache` interface defines methods for interacting with the second-level cache programmatically.

The **Cache** interface defines methods to check whether a particular entity has cached **data**, to remove a particular entity **from** the cache, to remove all instances (and instances of subclasses) of an entity **class from** the cache, and to clear the cache of all entity **data**.

Note - If the second-level cache has been disabled, calls to the **Cache** interface's methods have no effect, except for **contains**, which will always return **false**.

Checking Whether An Entity's Data is Cached

Call the **Cache.contains** method to find out whether a given entity is currently in the second-level cache.

The **contains** method returns **true** if the entity's **data** is cached, and **false** if the **data** is not in the cache.

```
EntityManager em = ...;  
Cache cache =  
em.getEntityManagerFactory().  
getCache();  
String personPK = ...;  
if (cache.contains  
(Person.class, personPK)) {  
    // the data is cached  
} else {  
    // the data is NOT cached  
}
```

Removing an Entity from the Cache

Call one of the **Cache.evict** methods to remove a particular entity or all entities of a given type **from** the second-level cache.

To remove a particular entity **from** the cache, call the **evict** method and pass in the entity **class** and the primary key of the entity:

```
EntityManager em = ...;  
Cache cache =  
em.getEntityManagerFactory().  
getCache();  
String personPK = ...;  
cache.evict(Person.class,  
personPK);
```

To remove all instances of a particular entity class, including subclasses, call the **evict** method and specify the entity class:

```
EntityManager em = ...;  
Cache cache =  
em.getEntityManagerFactory().  
getCache();  
cache.evict(Person.class);
```

All instances of the **Person** entity **class** will be removed **from** the cache.

If the **Person** entity has a subclass, **Student**, calls to the above method will remove all instances of **Student from** the cache as well.

Removing All Data from the Cache

Call the **Cache.evictAll** method to completely clear the second-level cache:

```
EntityManager em = ...;  
Cache cache =  
em.getEntityManagerFactory().  
getCache();  
cache.evictAll();
```