

Resource Connections

Java EE components can access a wide variety of resources, including **databases**, mail sessions, Java Message Service **objects**, and URLs.

The Java EE 6 platform provides mechanisms that allow you to access all these resources in a similar manner.

This chapter explains how to get connections to several types of resources. The following topics are addressed here:

- **Resources and JNDI Naming**
- **DataSource Objects and Connection Pools**
- **Resource Injection**
- **Resource Adapters and Contracts**
- **Metadata Annotations**
- **Common Client Interface**
- **Further Information about Resources**

Resources and JNDI Naming

In a distributed application, components need to access other components and resources, such as databases.

For example, a servlet might invoke remote methods on an enterprise bean that retrieves information from a database.

In the Java EE platform, the Java Naming and Directory **Interface (JNDI)** naming service enables components to locate other components and resources.

A **resource** is a program **object** that provides connections to **systems**, such as **database** servers and messaging systems.

(A Java **Database** Connectivity resource is sometimes referred to as a **data** source.)

Each resource **object** is identified by a unique, people-friendly name, called the JNDI name.

For example, the JNDI name of the **JDBC** resource for the Java DB **database** that is shipped with the GlassFish Server is **jdbc/___default.**

An administrator creates resources in a JNDI namespace.

In the GlassFish Server, you can use either the Administration Console or the `asadmin` command to create resources.

Applications then use annotations to inject the resources.

If an application uses resource injection, the GlassFish Server invokes the JNDI API, and the application is not required to do so.

However, it is also possible for an application to locate resources by making direct calls to the JNDI API.

A resource object and its JNDI name are bound together by the naming and directory service.

To create a **new** resource, a **new** name/object binding is entered **into** the JNDI namespace.

You inject resources by using the **@Resource** annotation in an application.

You can use a deployment **descriptor** to override the resource mapping that you **specify** in an annotation.

Using a deployment **descriptor** allows you to change an application by repackaging it rather than by both recompiling the source files and repackaging.

However, for most applications, a deployment **descriptor** is not necessary.

DataSource Objects and Connection Pools

To store, organize, and retrieve data, most applications use a relational database.

Java EE 6 components may access relational databases through the JDBC API.

For information on this API, see

<http://www.oracle.com/technetwork/java/javase/tech/index-jsp-136101.html>.

In the **JDBC API**, databases are accessed by using **DataSource** objects.

A **DataSource** has a set of properties that identify and describe the real-world **data** source that it represents.

These properties include such information as the location of the **database** server, the name of the **database**, the network protocol to use to communicate with the server, and so on.

In the GlassFish Server, a **data** source is called a **JDBC** resource.

Applications access a **data** source by using a connection, and a **DataSource** object can be thought of as a factory for connections to the particular **data** source that the **DataSource** instance represents.

In a basic **DataSource** implementation, a call to the **getConnection** method returns a connection **object** that is a physical connection to the **data** source.

A **DataSource** object may be registered with a JNDI naming service.

If so, an application can use the JNDI **API** to access that **DataSource** object, which can then be used to connect to the **data** source it represents.

DataSource objects that implement connection pooling also produce a connection to the particular **data** source that the **DataSource** class represents.

The connection **object** that the **getConnection** method returns is a handle to a **PooledConnection** **object** rather than being a physical connection.

An application uses the connection **object** in the same way that it uses a connection.

Connection **pooling** has no effect on application code except that a **pooled connection**, like all connections, should always be explicitly closed.

When an application closes a connection that is **pooled**, the connection is returned to a **pool** of reusable connections.

The next time **getConnection** is called, a handle to one of these **pooled** connections will be returned if one is available.

Because connection **pooling** avoids creating a **new** physical connection every time one is requested, applications can run significantly faster.

A **JDBC** connection pool is a group of reusable connections for a particular database.

Because creating each new physical connection is time consuming, the server maintains a pool of available connections to increase performance.

When it requests a connection, an application obtains one from the pool.

When an application closes a connection, the connection is returned to the pool.

Applications that use the Persistence API specify the DataSource object they are using in the jta-data-source element of the persistence.xml file:

```
<jta-data-source>  
jdbc/MyOrderDB  
</jta-data-source>
```

This is typically the only reference to a **JDBC object** for a persistence unit.

The application code does not refer to any **JDBC objects**.

Resource Injection

The `javax.annotation.Resource` annotation is used to declare a reference to a resource; `@Resource` can decorate a class, a field, or a method.

The container will inject the resource referred to by **@Resource** into the component either at runtime or when the component is initialized, depending on whether field/method injection or class injection is used.

With field-based and method-based injection, the container will inject the resource when the application is initialized.

For **class**-based injection, the resource is looked up by the application at runtime.

The **@Resource** annotation has the following elements:

- . **name**: The JNDI name of the resource
- . **type**: The Java language type of the resource
- . **authenticationType**: The authentication type to use for the resource

- **shareable**: Indicates whether the resource can be shared
- **mappedName**: A **nonportable**, implementation-specific name to which the resource should be mapped
- **description**: The **description** of the resource

The **name** element is the JNDI name of the resource and is optional for field-based and method-based injection.

For field-based injection, the default **name** is the field name qualified by the **class** name.

For method-based injection, the default **name** is the JavaBeans property name, based on the method qualified by the **class** name.

The **name** element must be specified for **class**-based injection.

The type of resource is determined by one of the following:

- . The type of the field the **@Resource** annotation is decorating for field-based injection
- . The type of the JavaBeans property the **@Resource** annotation is decorating for method-based injection
- . The **type** element of **@Resource**

For **class**-based injection, the **type** element is required.

The **authenticationType** element is used only for connection factory resources, such as the resources of a connector, also called the resource adapter, or **data** source.

This element can be set to one of the **javax.annotation.Resource**.

AuthenticationType enumerated type values: **CONTAINER**, the default, and **APPLICATION**.

The **shareable** element is used only for **Object Resource Broker (ORB)** instance resources or connection factory resource.

This element indicates whether the resource can be shared between this component and other components and may be set to **true**, the default, or **false**.

The **mappedName** element is a nonportable, implementation-specific name to which the resource should be mapped.

Because the **name** element, when **specified** or **defaulted**, is local only to the application, many Java EE servers provide a way of referring to resources across the application server.

This is done by setting the **mappedName** element.

Use of the **mappedName** element is **nonportable** across Java EE server implementations.

The **description** element is the **description** of the resource, typically in the default language of the system on which the application is deployed.

This element is used to help identify resources and to help application developers choose the correct resource.

Field-Based Injection

To use field-based resource injection, declare a field and decorate it with the `@Resource` annotation.

The container will infer the name and type of the resource if the `name` and `type` elements are not specified.

If you do specify the **type** element, it must match the field's **type** declaration.

In the following code, the container infers the **name** of the resource, based on the **class** name and the field name:

```
com.example.SomeClass/myDB.
```

The inferred **type** is

javax.sql.DataSource.class:

package com.example;

```
public class SomeClass {  
    @Resource  
    private javax.sql.DataSource myDB;  
    . . .  
}
```

In the following code, the JNDI name is `customerDB`, and the inferred type is `javax.sql.DataSource`:

```
package com.example;  
public class SomeClass {  
    @Resource(name="customerDB")  
    private javax.sql.DataSource myDB;  
    . . .  
}
```

Method-Based Injection

To use method-based injection, declare a setter method and decorate it with the `@Resource` annotation.

The container will infer the name and type of the resource if the `name` and `type` elements are not specified.

The setter method must follow the JavaBeans conventions for property names: The method name must begin with **set**, have a **void** return type, and only one parameter.

If you do **specify** the **type** element, it must match the field's type declaration.

In the following code, the container infers the **name** of the resource based on the **class** name and the field name:

`com.example.SomeClass/myDB.`

The inferred **type** is

`javax.sql.DataSource.class:`

```
package com.example;  
public class SomeClass {
```

```
private javax.sql.DataSource myDB;  
...  
@Resource  
private void  
setMyDB(javax.sql.DataSource ds)  
{ myDB = ds; }  
...  
}
```

In the following code, the JNDI name is `customerDB`, and the inferred type is `javax.sql.DataSource` class:

```
package com.example;  
public class SomeClass {  
    private javax.sql.DataSource myDB;  
    ...  
    @Resource(name="customerDB")  
    private void  
    setMyDB(javax.sql.DataSource ds)
```



```
{ myDB = ds; }  
  
...  
}
```

Class-Based Injection

To use **class**-based injection, decorate the **class** with a **@Resource** annotation, and set the required **name** and **type** elements:

```
@Resource (name="myMessageQueue",  
type="javax.jms.ConnectionFactory")  
public class SomeMessageBean { ... }
```

The **@Resources** annotation is used to group together multiple **@Resource** declarations for **class**-based injection.

The following code shows the `@Resources` annotation containing two `@Resource` declarations.

One is a Java Message Service message queue, and the other is a JavaMail session:

```
@Resources ( {  
@Resource (name="myMessageQueue",
```

```
type="javax.jms.ConnectionFactory")  
/  
@Resource(name="myMailSession",  
type="javax.mail.Session")  
})  
public class SomeMessageBean  
{...}
```

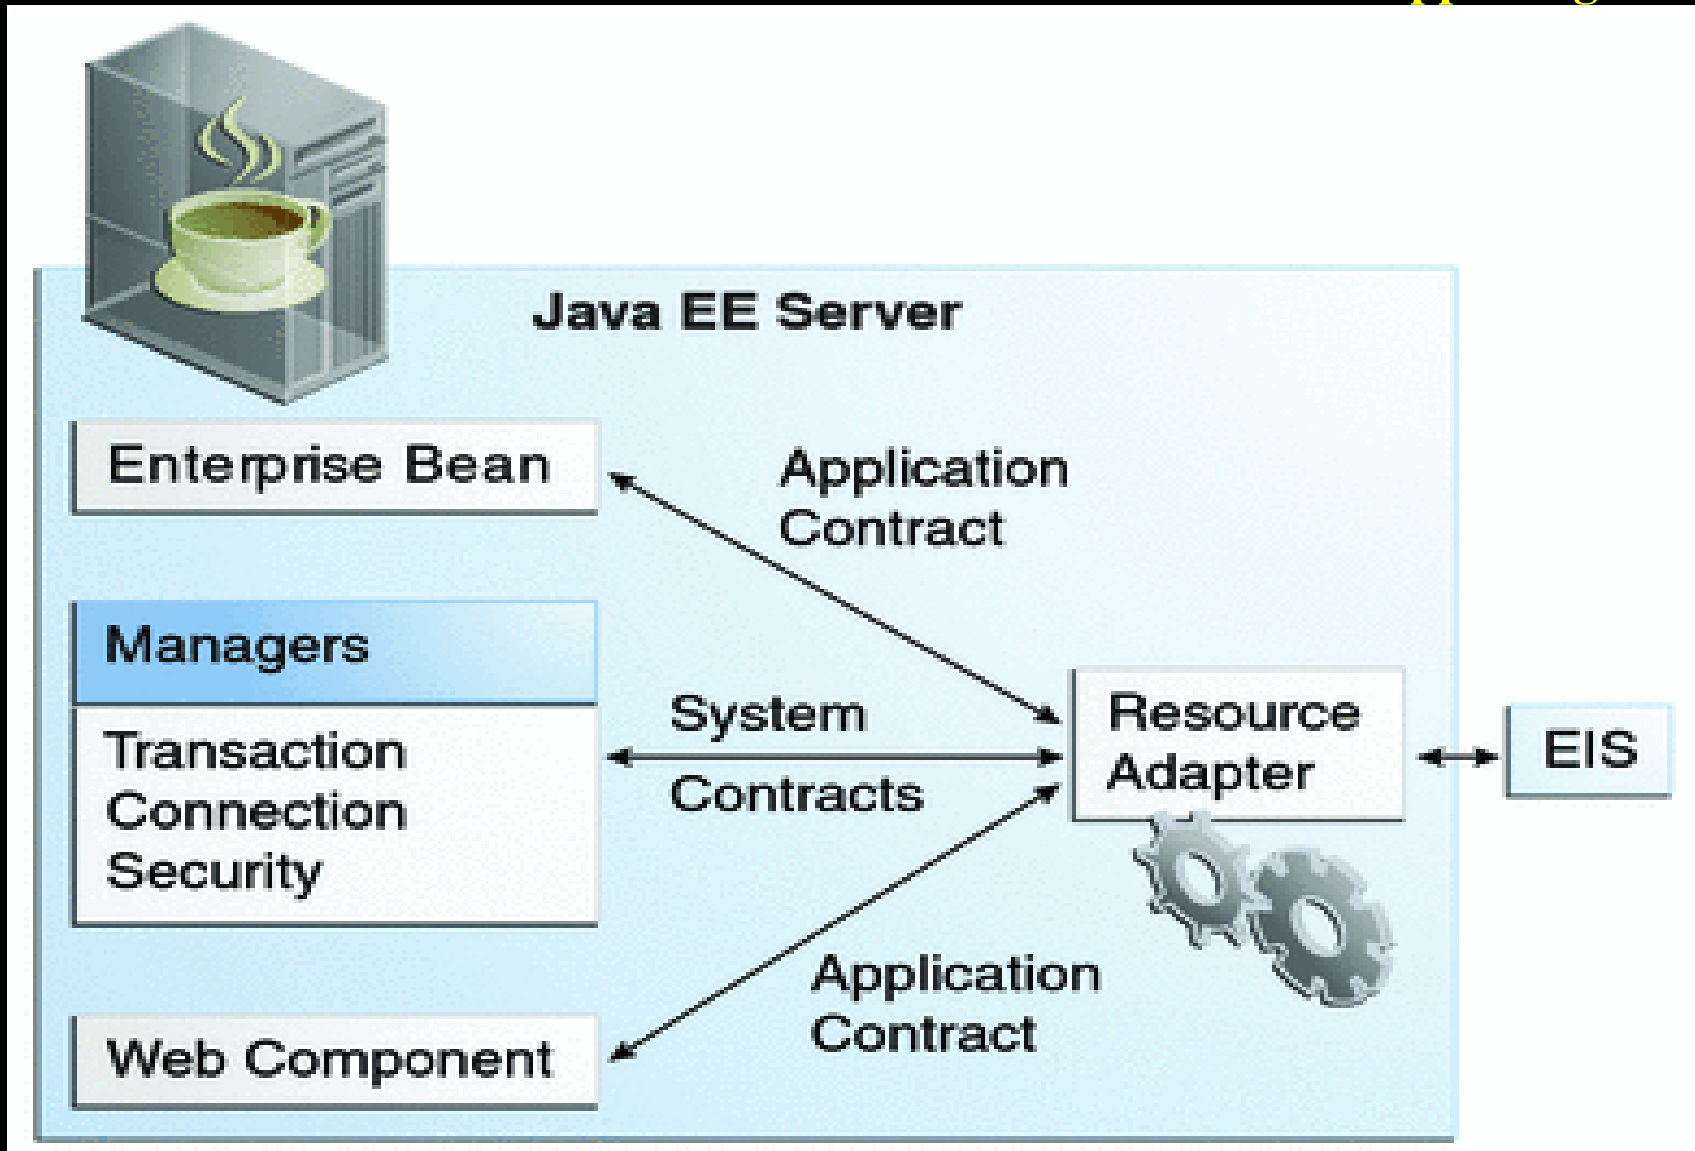
Resource Adapters and Contracts

A resource adapter is a Java EE component that implements the Java EE Connector **Architecture** for a **specific EIS**.

Examples of EISs include enterprise resource planning, mainframe transaction **processing**, and **database systems**.

As illustrated in Figure 44-1, the resource adapter facilitates communication between a Java EE application and an EIS.

Figure 44-1 Resource Adapters



Stored in a Resource Adapter Archive (RAR) file, a resource adapter can be deployed on any Java EE server, much like a Java EE application.

A RAR file may be contained in an Enterprise Archive (EAR) file, or it may exist as a separate file.

A resource adapter is analogous to a JDBC driver.

Both provide a standard **API** through which an application can access a resource that is outside the Java EE server.

For a resource adapter, the target system is an **EIS**; for a **JDBC** driver, it is a **DBMS**.

Resource adapters and **JDBC** drivers are rarely created by application developers.

In most cases, both types of software are built by vendors that sell tools, servers, or integration software.

The resource adapter mediates communication between the Java EE server and the EIS by means of contracts.

The application contract defines the **API** through which a Java EE component, such as an enterprise **bean**, accesses the EIS.

This **API** is the only view that the component has of the EIS.

The **system** contracts link the resource adapter to important services that are **managed** by the Java EE server.

The resource adapter itself and its system contracts are transparent to the Java EE component.

Management Contracts

The Java EE Connector **Architecture** defines system contracts that enable resource adapter lifecycle and thread **management**.

Lifecycle Management

The Connector **Architecture** specifies a lifecycle management contract that allows an application server to manage the lifecycle of a resource adapter.

This contract provides a mechanism for the application server to bootstrap a resource adapter instance during the deployment or application server startup.

This contract also provides a means for the application server to notify the resource adapter instance when it is undeployed or when an orderly shutdown of the application server takes place.

Work Management Contract

The Connector **Architecture** work **management** contract ensures that resource adapters use threads in the proper, recommended manner.

This contract also enables an application server to **manage** threads for resource adapters.

Resource adapters that improperly use threads can jeopardize the entire application server environment.

For example, a resource adapter might create too many threads or might not properly release threads it has created.

Poor thread handling inhibits application server shutdown and impacts the application server's performance because creating and destroying threads are expensive operations.

The work **management** contract establishes a means for the application server to **pool** and **reuse** threads, similar to **pooling** and reusing connections.

By adhering to this contract, the resource adapter does not have to **manage** threads itself.

Instead, the resource adapter has the application server create and provide needed threads.

When it is finished with a given thread, the resource adapter returns the thread to the application server.

The application server manages the thread, either returning it to a pool for later reuse or destroying it.

Handling threads in this manner results in increased application server performance and more efficient use of resources.

In addition to moving thread management to the application server, the Connector Architecture provides a flexible model for a resource adapter that uses threads.

- . The requesting thread can choose to block (stop its own execution) until the work thread completes.

- . The requesting thread can block while it waits to get the work thread.**

When the application server provides a work thread, the requesting thread and the work thread execute in parallel.

- . The resource adapter can opt to submit the work for the thread to a queue.**

The thread executes the work **from** the queue at some later **point**.

The resource adapter continues its own execution **from** the **point** it submitted the work to the queue, no matter when the thread executes it.

With the latter two approaches, the submitting thread and the work thread may execute simultaneously or independently.

For these approaches, the contract specifies a listener mechanism to notify the resource adapter that the thread has completed its operation.

The resource adapter can also specify the execution context for the thread, and the work management contract controls the context in which the thread executes.

Generic Work Context Contract

The work **management** contract between the application server and a resource adapter enables a resource adapter to do a task, such as communicating with the EIS or delivering messages, by delivering **Work** instances for execution.

A generic work context contract enables a resource adapter to control the contexts in which the **Work** instances that it submits are executed by the application server's **WorkManager**.

A generic work context mechanism also enables an application server to support **new** message inflow and delivery schemes.

It also provides a richer contextual **Work** execution environment to the resource adapter while still **maintaining** control over concurrent **behavior** in a **managed** environment.

The generic work context contract standardizes the transaction context and the security context.

Outbound and Inbound Contracts

The Connector **Architecture** defines the following outbound contracts, system-level contracts between an application server and an EIS that enable outbound connectivity to an EIS.

- . The connection management contract supports connection pooling, a technique that enhances application performance and scalability.

Connection pooling is transparent to the application, which simply obtains a connection to the EIS.

- . The transaction management contract extends the connection management contract and provides support for management of both local and XA transactions.

A local transaction is limited in scope to a single EIS system, and the EIS resource manager itself manages such transaction.

An XA transaction or global transaction can span multiple resource managers.

This form of transaction requires transaction coordination by an external transaction manager, typically bundled with an application server.

A transaction manager uses a two-phase commit protocol to manage a transaction that spans multiple resource managers or EISs, and uses one-phase commit optimization if only one resource manager is participating in an XA transaction.

- The security **management** contract provides mechanisms for authentication, authorization, and secure communication between a Java EE server and an EIS to protect the information in the EIS.

A work security map matches EIS identities to the application server domain's identities.

Inbound contracts are system contracts between a Java EE server and an EIS that enable inbound connectivity from the EIS: pluggability contracts for message providers and contracts for importing transactions.

Metadata Annotations

Java EE Connector **Architecture** 1.6 **introduces** a set of annotations to minimize the need for deployment **descriptors**.

- . The **@Connector** annotation can be used by the resource adapter developer to specify that the **JavaBeans** component is a resource adapter **JavaBeans** component.

This annotation is used for providing metadata about the capabilities of the resource adapter.

Optionally, you can provide a **JavaBeans** component implementing the **ResourceAdapter** interface, as in the following example:

```
@Connector(  
description =  
"Sample adapter using the JavaMail API",  
displayName =  
"InboundResourceAdapter",  
vendorName = "My Company, Inc.",
```

```
eisType = "MAIL",  
version = "1.0"  
)  
  
public class ResourceAdapterImpl  
implements ResourceAdapter,  
java.io.Serializable {...}
```

- The **@ConnectionDefinition** annotation defines a set of connection **interfaces** and **classes** pertaining to a particular connection type, as in the following example:

```
@ConnectionDefinition (  
    connectionFactory =  
        samples.mailra..api.  
        JavaMailConnectionFactory.class,  
    connectionFactoryImpl =  
        samples.mailra.ra.outbound.  
        JavaMailConnectionFactoryImpl.class  
/  
    connection =  
        samples.connectors.mailconnector.  
        api.JavaMailConnection.class,
```

```
connectionImpl =
samples.mailra..ra.outbound.
JavaMailConnectionImpl.class
)
public class
ManagedConnectionFactoryImpl
implements
ManagedConnectionFactory,
Serializable {... ..
@ConfigProperty
(defaultValue = "UnknownHostName")
```



```
public void setServerName  
(String serverName) { ... }  
}
```

- . The `@AdministeredObject` annotation designates a JavaBeans component as an administered object.

- The `@Activation` annotation contains configuration information pertaining to inbound connectivity **from** an EIS instance, as in the following example:

```
@Activation(  
messageListeners = {  
samples.mailra.api.  
JavaMailMessageListener.class  
} )
```

```
public class ActivationSpecImpl
implements
javax.resource.spi.ActivationSpec,
java.io.Serializable {...
@ConfigProperty()
// serverName property value
private String serverName =
new String("");
@ConfigProperty()
// userName property value
```

```
private String userName =  
new String("");  
@ConfigProperty()  
// password property value  
private String password =  
new String("");  
@ConfigProperty()  
// folderName property value  
private String folderName =  
new String("Inbox");  
// protocol property value
```

```
// Normally imap or pop3
@ConfigProperty(
    description =
        "Normally imap or pop3"
)
private String protocol =
    new String("imap");
...
...
}
```

- The `@ConfigProperty` annotation can be used on JavaBeans components to provide additional configuration information that may be used by the deployer and resource adapter provider.

The preceding example code shows several `@ConfigProperty` annotations.

The specification allows a resource adapter to be developed in mixed-mode form, that is the ability for a resource adapter developer to use both metadata annotations and deployment descriptors in applications.

An application assembler or deployer may use the deployment descriptor to override the metadata annotations specified by the resource adapter developer.

The deployment descriptor for a resource adapter is named **ra.xml**.

The **metadata-complete** attribute defines whether the deployment descriptor for the resource adapter module is complete or whether the **class** files available to the module and packaged with the resource adapter need to be examined for annotations that **specify** deployment information.

For the complete list of annotations and JavaBeans components introduced in the Java EE 6 platform, see the Java EE Connector Architecture 1.6 specification.

Common Client Interface

This section explains how components use the Connector **Architecture** Common Client Interface (CCI) **API** and a resource adapter to access **data** from an EIS.

The CCI **API** defines a set of **interfaces** and **classes** whose methods allow a client to perform typical **data** access operations.

The CCI interfaces and classes are as follows:

- . **ConnectionFactory**: Provides an application component with a **Connection** instance to an EIS.
- . **Connection**: Represents the connection to the underlying EIS.

- **ConnectionSpec**: Provides a means for an application component to pass connection-request-specific properties to the **ConnectionFactory** when making a connection request.
- **Interaction**: Provides a means for an application component to execute EIS functions, such as **database** stored procedures.

- **InteractionSpec**: Holds properties pertaining to an application component's interaction with an EIS.
- **Record**: The superinterface for the various kinds of record instances.

Record instances can be **MappedRecord**, **IndexedRecord**, or **ResultSet** instances, all of which inherit from the **Record** interface.

- **RecordFactory**: Provides an application component with a **Record** instance.
- **IndexedRecord**: Represents an ordered collection of **Record** instances based on the **java.util.List** interface.

A client or application component that **uses** the CCI to **interact** with an underlying EIS does so in a prescribed manner.

The component must establish a connection to the EIS's resource **manager**, and it does so using the **ConnectionFactory**.

The **Connection** **object** represents the connection to the EIS and is **used** for subsequent **interactions** with the EIS.

The component performs its interactions with the EIS, such as accessing data from a specific table, using an Interaction object.

The application component defines the Interaction object by using an InteractionSpec object.

When it reads **data from** the EIS, such as **from database tables**, or writes to those **tables**, the application component does so by using a particular type of **Record** instance: a **MappedRecord**, an **IndexedRecord**, or a **ResultSet** instance.

Note, too, that a client application that relies on a CCI resource adapter is very much like any other Java EE client that uses enterprise **bean** methods.

Further Information about Resources

For more information about resources and annotations, see

- Java EE 6 Platform Specification (JSR 316):

<http://jcp.org/en/jsr/detail?id=316>

- Java EE Connector Architecture 1.6 specification:

<http://jcp.org/en/jsr/detail?id=322>

- . **EJB 3.1 specification:**

<http://jcp.org/en/jsr/detail?id=318>

- . **Common Annotations for the Java Platform:**

<http://www.jcp.org/en/jsr/detail?id=250>